Jorge Izaguirre García

Forward Deployed AI Engineer | Agentic Systems & MCP | LLM Evals | Full-Stack + AWS

Personal data

Address:	Las Rozas de Madrid, Spain (open to travel)
Phone number:	+34 684 068 786
Email address:	jorge@izaguirre.eu

Professional summary

Computer Engineer (UPM, thesis 9.6/10) who takes AI systems from client discovery to production. In 12 months, delivered AI and full-stack solutions for 5 clients (legal, compliance, ticketing, e-commerce, healthcare education). Lead architect of an AI email-processing pipeline handling ∼8,400 emails/month that let a law firm go from 6 to 2 people on the process, and top backend contributor to an agentic compliance SaaS with a production MCP server. Works directly with client stakeholders through demos, technical proposals, go-live and iteration. Native Spanish, English C2 (Cambridge Proficiency).

Technical experience

AI Engineer at Galileo Studio. 07/2026–Present (AI Engineering Intern 02/2026–07/2026)

Pyramid — law firm managing traffic fines. Lead engineer & architect, 03/2026–Present.

- Designed and built an AI pipeline that classifies inbound emails, extracts fine data (OCR + schema-constrained LLM output) and files it into the client's systems via their APIs, Microsoft Graph and Power Automate. Python, AWS Lambda/SQS/DynamoDB/S3/Textract, OpenAI, Terraform, CodePipeline.
- ∼8,400 emails/month in production at a 99% success rate. Raised fully hands-off automation from 0% to ∼57% in 4 months by mining production KPIs for automatable clusters; handling time cut from ∼10 to 2–3 min per email, taking the client's team from 6 to 2 people (∼1,000 h/month saved).
- Built the evaluation process (human labeling + LLM-as-judge): 97.4% / 90.4% classification precision on validated production emails, ∼800 automated tests, query-answer accuracy raised from ∼25% to ∼75–80%.
- Led the model migration ahead of OpenAI's fine-tuning deprecation: offline evaluation of 11 candidates (OpenAI, Anthropic, open-source), then fine-tuned Qwen3.8 27B (LoRA, Together AI) to near-parity on replayed production traffic; caught training-data memorization through functional testing before rollout. Presented the technical proposal (SageMaker deployment) to the client.
- Owned production: traced recurring CloudWatch failures to a payload limit in the client's API, aligned with their tech team, measured quality loss and shipped 10 MB→1 MB document compression with no measurable degradation. Runs a nightly reviewer agent that audits code and logs and files issues.

Kura — Galileo's compliance-audit SaaS (alpha). Top backend contributor, 09/2026–Present.

- Re-architected a hardcoded single-framework app (criminal compliance) into a multi-vertical platform driven by declarative YAML: data model, API, schema-generated UI, auto-generated MCP tools and CI compatibility gates, shipped in ∼7 working days. Built a POC first to win buy-in from non-technical leadership.
- Built an AI interview-driven skill that generates new verticals in ∼10 min with no code or migrations; used to add China MLPS (myself), ISO 27001 and France Sapin II (teammates).
- Designed and shipped to production an MCP server letting any LLM agent read and write audit data: Clerk OAuth, multi-tenant Postgres RLS, draft→confirm writes, prompt-injection safeguards and auditing of agent actions. Multi-provider LLM layer (Claude, OpenAI, Gemini).
- 41% of backend PRs in the last month (∼74 PRs in 6 weeks), in a CI pipeline with strict typing and a 688-test RLS security suite. Co-designed features with compliance experts and demoed the architecture to the technical team.

Other client deliveries

- **Acatia** (whistleblowing SaaS, ∼125 tenants): designed the AWS infrastructure-as-code (Terraform: Lambda/Bref, RDS Proxy, CloudFront, API Gateway, VPC) and the migration cost analysis; cut local API latency from ∼11 s to ∼1.5 s; built a multi-tenant superadmin portal.

- **Vibralo** (event ticketing): migrated the payment gateway (Paylands → Redsys); ran a security audit and fixed the findings (unauthenticated endpoints, RLS gaps, SSRF, leaked service key, double-refund race condition, overselling).
- **Fraeternity** (custom apparel e-commerce): Sendcloud shipping integration, queued payment notifications, address validation with suggestions (FastAPI, React, Fly.io).
- **WorkOps** (multi-agent Word assistant): LangGraph agent tools for document editing, page-level PDF citations, web search and response streaming.

Freelance Software Engineer — Rey Juan Carlos University. 11/2025–06/2026

- Built a web app for clinical tutors to evaluate nursing trainees (React / TypeScript, Node.js, SQL), iterating with medical stakeholders.

Projects

- **PR explainer videos** (in development): an agentic pipeline that turns each pull request (diff, Linear context, review comments) into a narrated video: LLM-written script, ElevenLabs TTS synced to motion graphics, rendered with ffmpeg.
- **Hackathon Winner** — Best Madrid UPM 2023 (4-person team).

Academic education

- 2022–2026. B.Sc. Computer Engineering. Universidad Politécnica de Madrid (UPM). Thesis (9.6/10): agentic document retrieval without vector databases (LangGraph).
- 2020–2022. Science and Technology Bachelor's Degree. Colegio Logos International School.

Technical skills

- AI / Agents: OpenAI, Anthropic Claude, MCP server design, tool calling, LangGraph, RAG, LLM evaluation (LLM-as-judge, offline multi-model), LoRA fine-tuning, prompt-injection mitigation, AWS Textract, ElevenLabs
- Languages: Python, TypeScript, JavaScript, SQL; also Java, C, C++, C#, PHP
- Backend / Data: FastAPI, Node.js, PostgreSQL (RLS, JSONB), DynamoDB, MySQL, Supabase, REST APIs, webhooks, queues
- Cloud / DevOps: AWS (Lambda, SQS, DynamoDB, S3, API Gateway, CloudFront, CloudWatch), Terraform, Docker, GitHub Actions, CodePipeline
- Frontend: React, Next.js, Vite
- Delivery: client discovery, demos, technical proposals, production debugging, AI-native development (coding agents with test and review gates)

Language skills

- Spanish, native.
- English, C2 (Cambridge Certificate of Proficiency in English).
- French, B2.